

Operating Systems (234123)

Processes & Signals

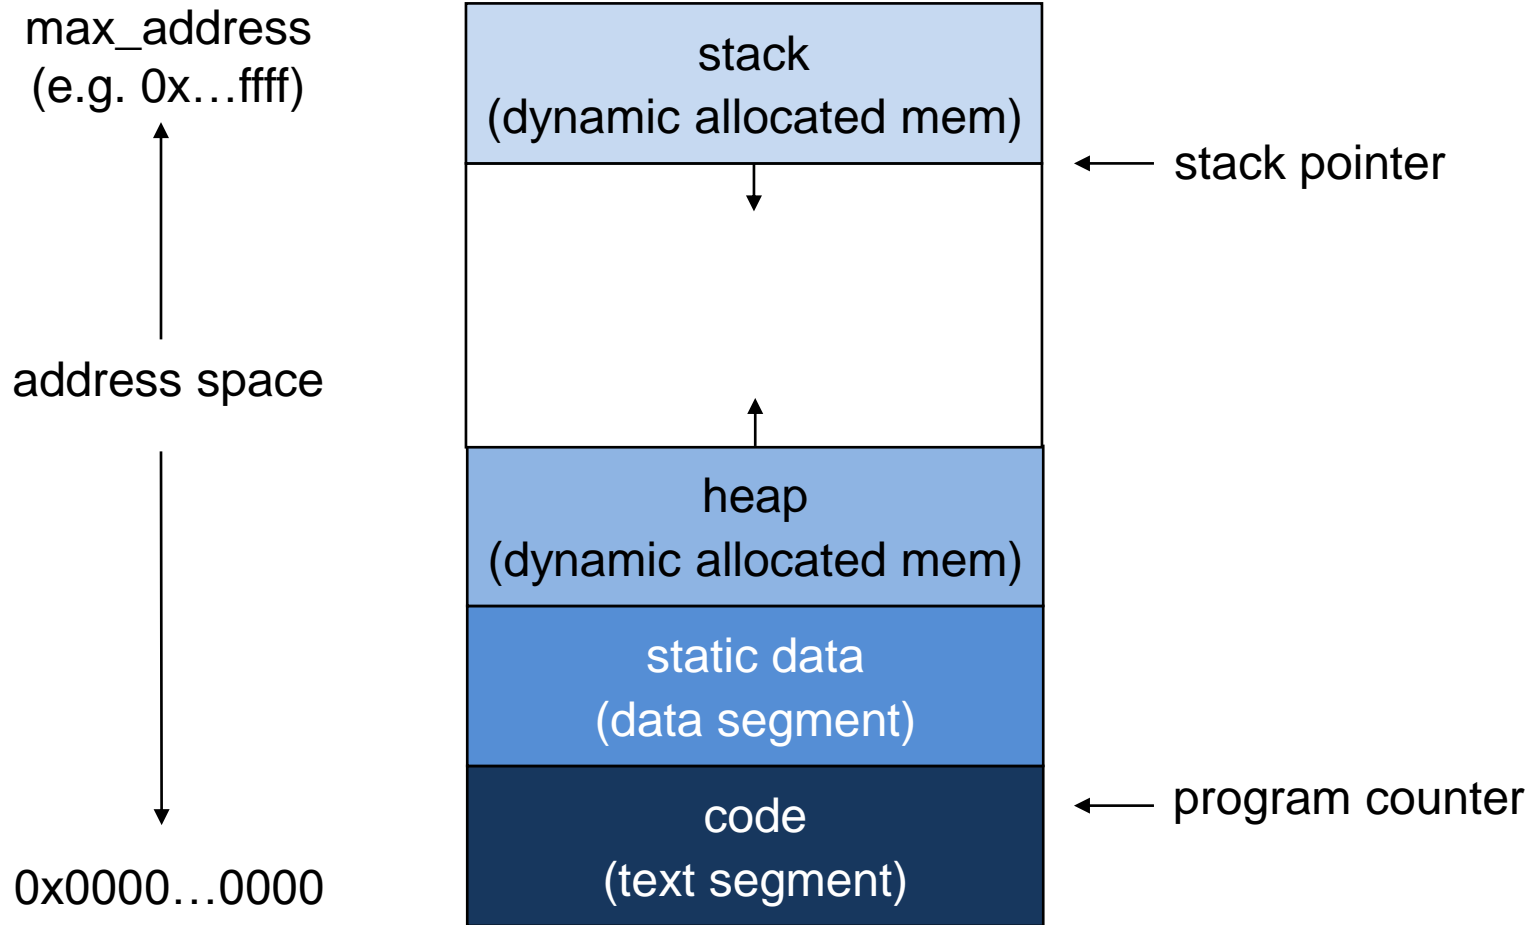
Dan Tsafir

2024-06-03, 2024-06-10

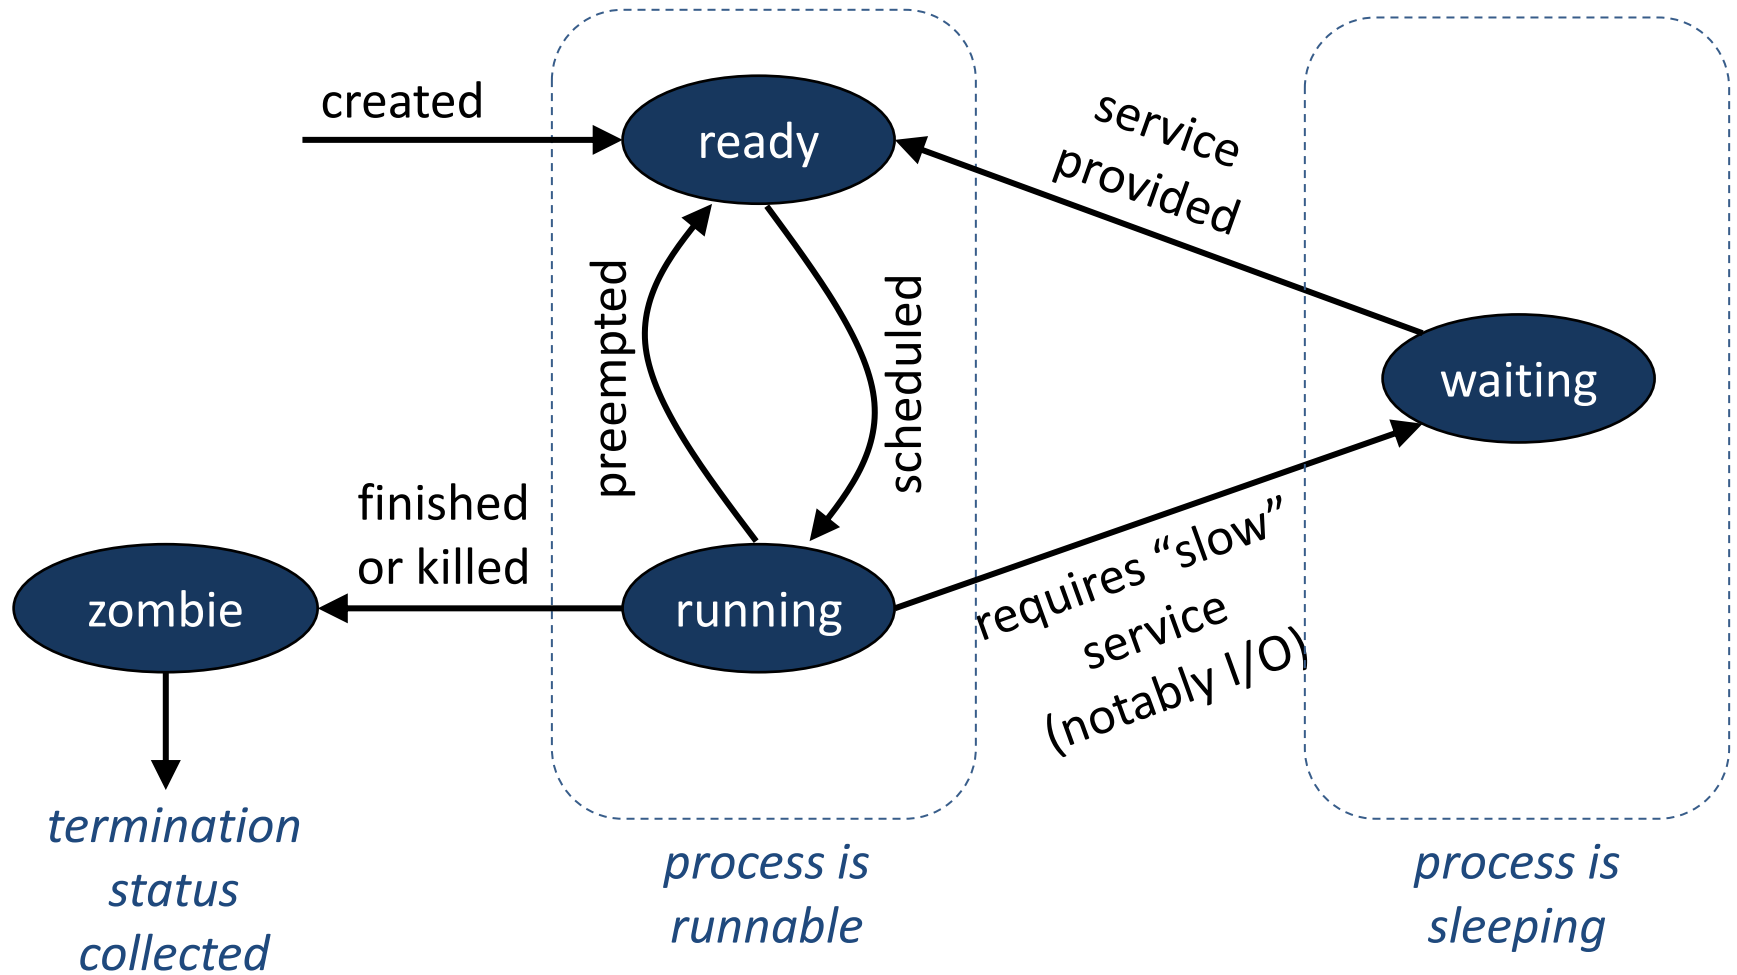
What's a process

- **An implementation of the abstract machine concept**
 - Which we discussed in the previous lecture
- **A running instance of an executable, invoked by a user**
 - Can have multiple independent processes of the same executable
- **A schedulable entity, on the CPU**
 - OS decides which of these entities gets to run on a CPU core, and when
- **Sometimes called**
 - Task or job
- **The OS kernel is neither a process nor a schedulable entity**
 - Rather, it's a set of procedures executing in response to events (\approx interrupts)
 - Albeit sometimes the OS runs some code within schedulable entities
 - But then we prefer not to refer to these entities as “processes”, which correspond to *user* programs; we may refer to them as “kernel threads” instead

Process address space is contiguous



Process states



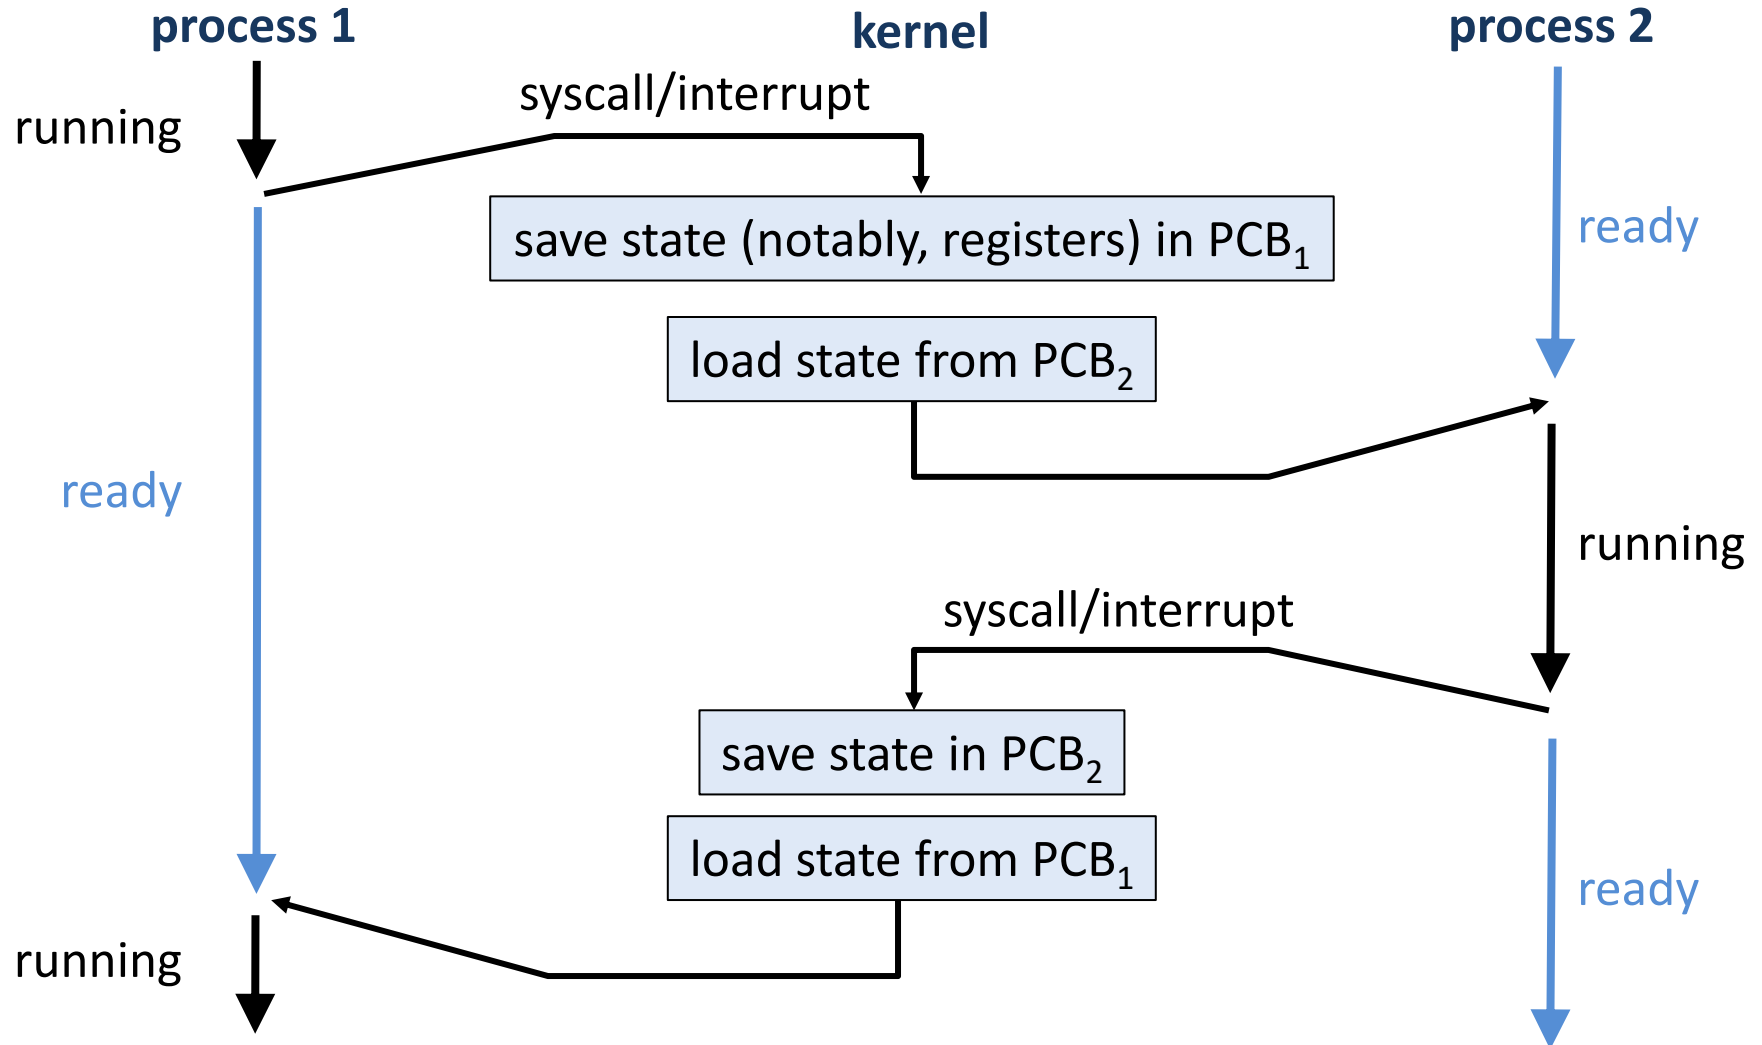
Process control block (“PCB”)

- **The OS maintains a “state” for every process**
 - Encapsulated in a PCB
- **In Linux**
 - Called a “process descriptor”
 - Of type `task_struct` (C struct)
 - Has $O(100)$ fields
- **Used in context switches**
 - Updated upon preemption
 - Loaded upon resumption
- **Question**
 - Can a process access its own PCB?

- PID (process ID)
- UID (user ID)
- Pointer to address space
- Registers
- Scheduling priority
- Resources usage limits (e.g., memory, CPU, num of open files)
- Resources consumed
- State (previous slide)
- Current/present working directory (`pwd=cwd`)
- Open files table
- ...

process attributes
in PCB

Context switching (in runnable state)



Process creation & termination

- **One process (the “parent”) can create another (the “child”)**
 - A new PCB is allocated and initialized
 - Homework: run ‘ps auxwww’ in the shell; PPID is the parent’s PID
- **In POSIX, child process inherits most of parent’s attributes**
 - UID, open files (should be closed if unneeded; why?), cwd, etc.
- **While executing, PCB moves between different queues**
 - According to state change graph
 - Queues: runnable, sleep/wait for event i ($i=1,2,3\dots$)
- **After a process dies (exit()s / interrupted), it becomes a zombie**
 - Parent uses `wait*` syscall to clear zombie from the system (why?)
 - Wait syscall family: `wait`, `waitpid`, `waitid`, `wait3`, `wait4`; example:
 - `pid_t wait4(pid_t, int *wstatus, int options, struct rusage *rusage);`
- **Parent can sleep/wait for its child to finish or run in parallel**
 - `wait*()` will block unless `WNOHANG` given in ‘options’
 - Homework: read ‘man 2 wait’

fork() – spawn a child process

- **fork() initializes a new PCB**
 - Based on parent's value
 - PCB added to runnable queue
- **Now there are 2 processes**
 - At same execution point
- **Child's new address space**
 - Complete copy of parent's space, with one difference...
- **fork() returns twice**
 - At the parent, with pid>0
 - At the child, with pid=0
- **What's the printing order?**
- **'errno' – a global variable**
 - Holds error num of last syscall

```
int main(int argc, char *argv[])
{
    int pid = fork();
    if( pid==0 ) {
        //
        // child
        //
        printf("parent=%d son=%d\n",
               getpid(), getpid());
    }
    else if( pid > 0 ) {
        //
        // parent
        //
        printf("parent=%d son=%d\n",
               getpid(), pid);
    }
    else { // print string associated
           // with errno
        perror("fork() failed");
    }
    return 0;
}
```


System call errors

```
// int errno = number of last system call error.
// Errors aren't zero. (If you want to test value of
// errno after a system call, need to zero it before.)
#include <errno.h> // see man 3 errno

// const char * const sys_errlist[];
// char* strerror(int errnum) {
//     // check errnum is in range
//     return sys_errlist[errnum];
// }
#include <string.h>

// void perror(const char *prefix);
// prints: "%s: %s\n" , prefix, sys_errlist[errno]
#include <stdio.h>
```

exec*() – replace current process image

- **To start an entirely new program**
 - Use the exec*() syscall family; for example:
 - int `execv`(const char *programPath, char *const argv[]);
 - Homework: read 'man execv'
- **Semantics**
 - Stops the execution of the invoking process
 - Loads the executable 'programPath'
 - Starts 'programPath', with 'argv' as its argv
 - Never returns (unless fails)
 - *Replaces* the new process; doesn't create a new process
 - In particular, PID and PPID are the same before/after exec*()