

Operating Systems (234123)

Processes & Signals

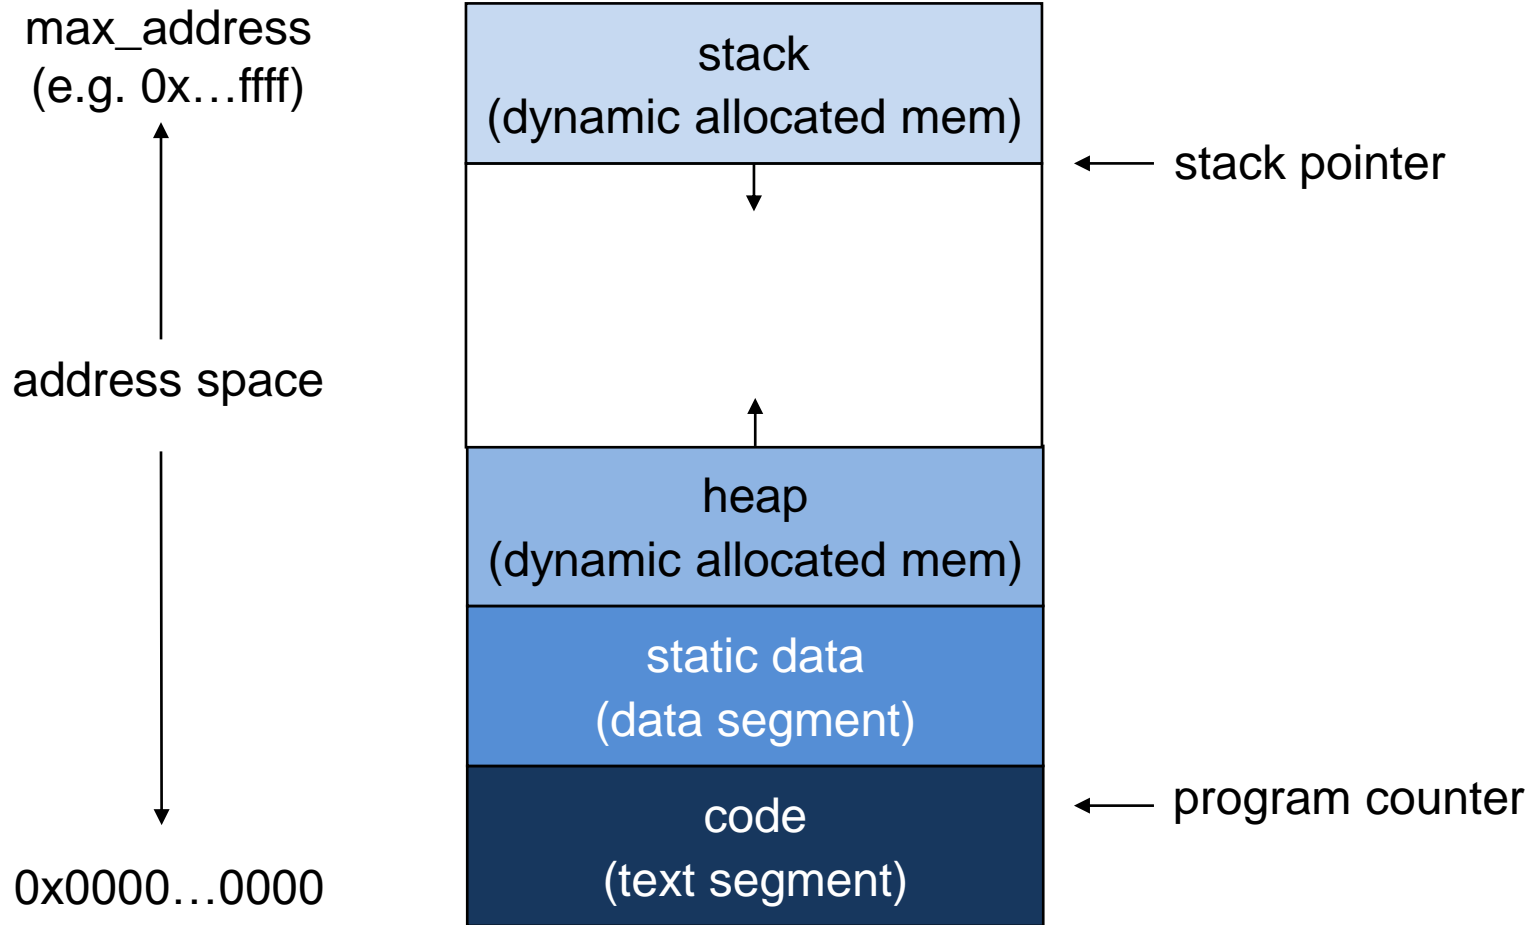
Dan Tsafir

2024-06-03, 2024-06-10

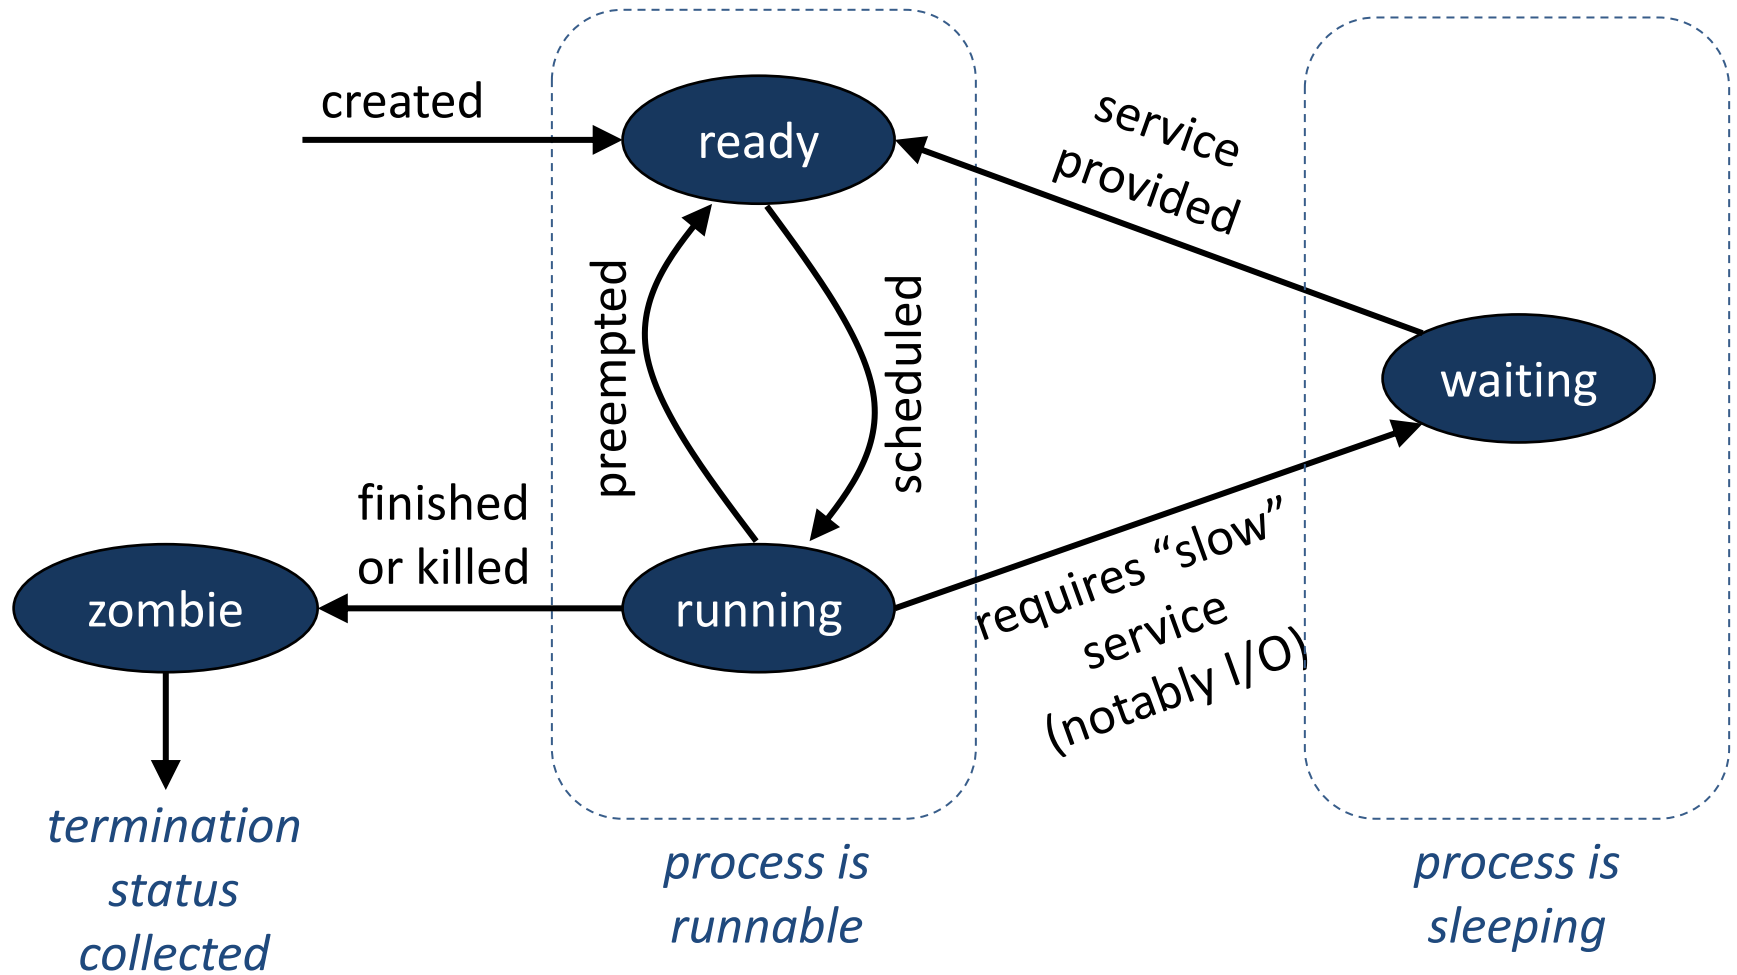
What's a process

- **An implementation of the abstract machine concept**
 - Which we discussed in the previous lecture
- **A running instance of an executable, invoked by a user**
 - Can have multiple independent processes of the same executable
- **A schedulable entity, on the CPU**
 - OS decides which of these entities gets to run on a CPU core, and when
- **Sometimes called**
 - Task or job
- **The OS kernel is neither a process nor a schedulable entity**
 - Rather, it's a set of procedures executing in response to events (\approx interrupts)
 - Albeit sometimes the OS runs some code within schedulable entities
 - But then we prefer not to refer to these entities as “processes”, which correspond to *user* programs; we may refer to them as “kernel threads” instead

Process address space is contiguous



Process states



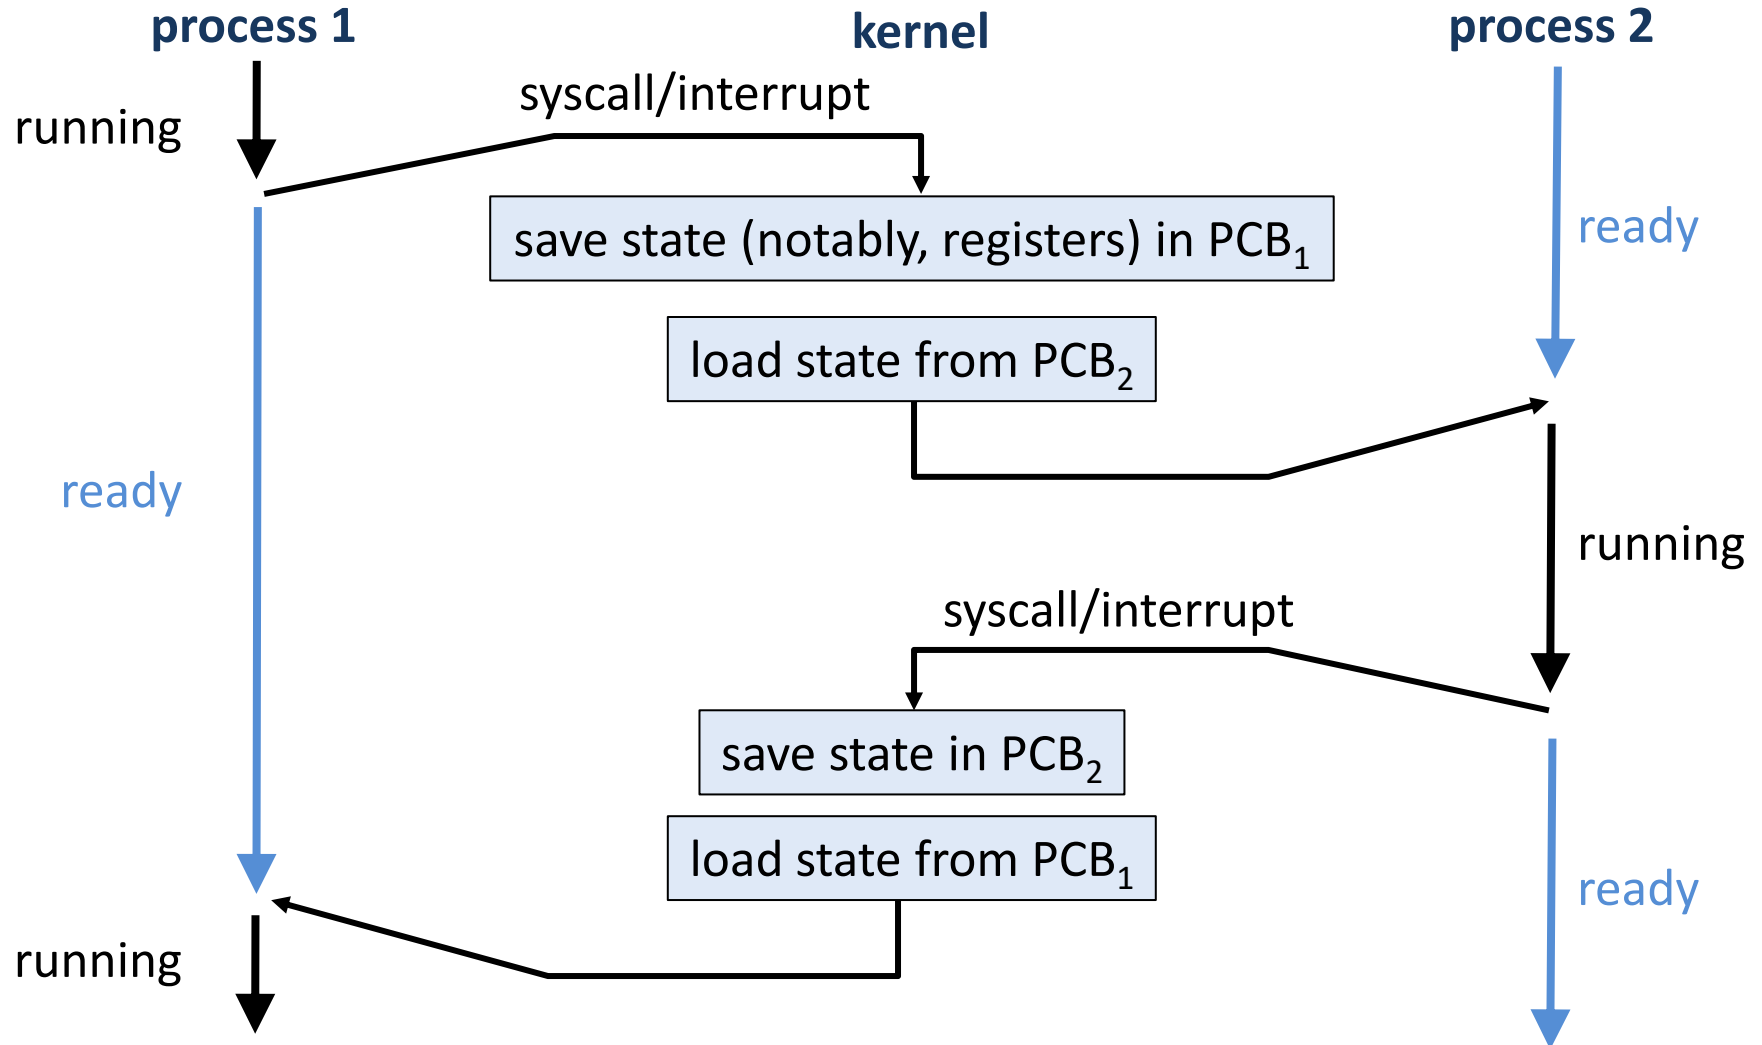
Process control block (“PCB”)

- **The OS maintains a “state” for every process**
 - Encapsulated in a PCB
- **In Linux**
 - Called a “process descriptor”
 - Of type `task_struct` (C struct)
 - Has $O(100)$ fields
- **Used in context switches**
 - Updated upon preemption
 - Loaded upon resumption
- **Question**
 - Can a process access its own PCB?

- PID (process ID)
- UID (user ID)
- Pointer to address space
- Registers
- Scheduling priority
- Resources usage limits (e.g., memory, CPU, num of open files)
- Resources consumed
- State (previous slide)
- Current/present working directory (`pwd=cwd`)
- Open files table
- ...

process attributes
in PCB

Context switching (in runnable state)



Process creation & termination

- **One process (the “parent”) can create another (the “child”)**
 - A new PCB is allocated and initialized
 - Homework: run ‘ps auxwww’ in the shell; PPID is the parent’s PID
- **In POSIX, child process inherits most of parent’s attributes**
 - UID, open files (should be closed if unneeded; why?), cwd, etc.
- **While executing, PCB moves between different queues**
 - According to state change graph
 - Queues: runnable, sleep/wait for event i ($i=1,2,3\dots$)
- **After a process dies (exit()s / interrupted), it becomes a zombie**
 - Parent uses `wait*` syscall to clear zombie from the system (why?)
 - Wait syscall family: `wait`, `waitpid`, `waitid`, `wait3`, `wait4`; example:
 - `pid_t wait4(pid_t, int *wstatus, int options, struct rusage *rusage);`
- **Parent can sleep/wait for its child to finish or run in parallel**
 - `wait*()` will block unless `WNOHANG` given in ‘options’
 - Homework: read ‘man 2 wait’

fork() – spawn a child process

- **fork() initializes a new PCB**
 - Based on parent's value
 - PCB added to runnable queue
- **Now there are 2 processes**
 - At same execution point
- **Child's new address space**
 - Complete copy of parent's space, with one difference...
- **fork() returns twice**
 - At the parent, with pid>0
 - At the child, with pid=0
- **What's the printing order?**
- **'errno' – a global variable**
 - Holds error num of last syscall

```
int main(int argc, char *argv[])
{
    int pid = fork();
    if( pid==0 ) {
        //
        // child
        //
        printf("parent=%d son=%d\n",
               getpid(), getpid());
    }
    else if( pid > 0 ) {
        //
        // parent
        //
        printf("parent=%d son=%d\n",
               getpid(), pid);
    }
    else { // print string associated
           // with errno
        perror("fork() failed");
    }
    return 0;
}
```


System call errors

```
// int errno = number of last system call error.
// Errors aren't zero. (If you want to test value of
// errno after a system call, need to zero it before.)
#include <errno.h> // see man 3 errno

// const char * const sys_errlist[];
// char* strerror(int errnum) {
//     // check errnum is in range
//     return sys_errlist[errnum];
// }
#include <string.h>

// void perror(const char *prefix);
// prints: "%s: %s\n" , prefix, sys_errlist[errno]
#include <stdio.h>
```

exec*() – replace current process image

- **To start an entirely new program**
 - Use the exec*() syscall family; for example:
 - int `execv`(const char *programPath, char *const argv[]);
 - Homework: read 'man execv'
- **Semantics**
 - Stops the execution of the invoking process
 - Loads the executable 'programPath'
 - Starts 'programPath', with 'argv' as its argv
 - Never returns (unless fails)
 - *Replaces* the new process; doesn't create a new process
 - In particular, PID and PPID are the same before/after exec*()

Simplistic UNIX shell loop example

Finished
here

```
int main(int argc, char *argv[])
{
    for(;;) {
        int stat;
        char **argv;
        char *c = readNextCom(&argv);
        int pid = fork();

        if( pid < 0 ) {
            perror("fork failed");
        }
        else if( pid==0 ) { // child
            execv(c, argv);
            perror("execv failed");
        }
        else { // parent
            if( wait(&stat) < 0 )
                perror("wait failed");
            else
                chkStatus(pid, stat);
            release(argv);
        }
    }
    return 0;
}
```

```
void chkStatus(int pid, int stat)
{
    if( WIFEXITED(stat) ) {
        printf("%d exit code=%d\n",
               pid, WEXITSTATUS(stat));
    }
    else if( WIFSIGNALED(stat) ) {
        // the topic we're going
        // to learn next
        printf("%d died on signal=%d\n",
               pid, WTERMSIG(stat));
    }
    else if
        // a few more options...
}
```

Who wait()-s for an “orphan” process?

- **POSIX specification says:**

- *“If a parent process terminates without waiting for all of its child processes to terminate, the remaining child processes shall be assigned a new parent process ID corresponding to an implementation-defined system process”*

- <https://pubs.opengroup.org/onlinepubs/9699919799/functions/wait.html>

- **Linux manual says:**

- *“init [...becomes] the parent of all processes whose natural parents have died, and it is responsible for reaping those when they die. [...] init expects to have a process id of 1”*

- <https://linux.die.net/man/8/init>

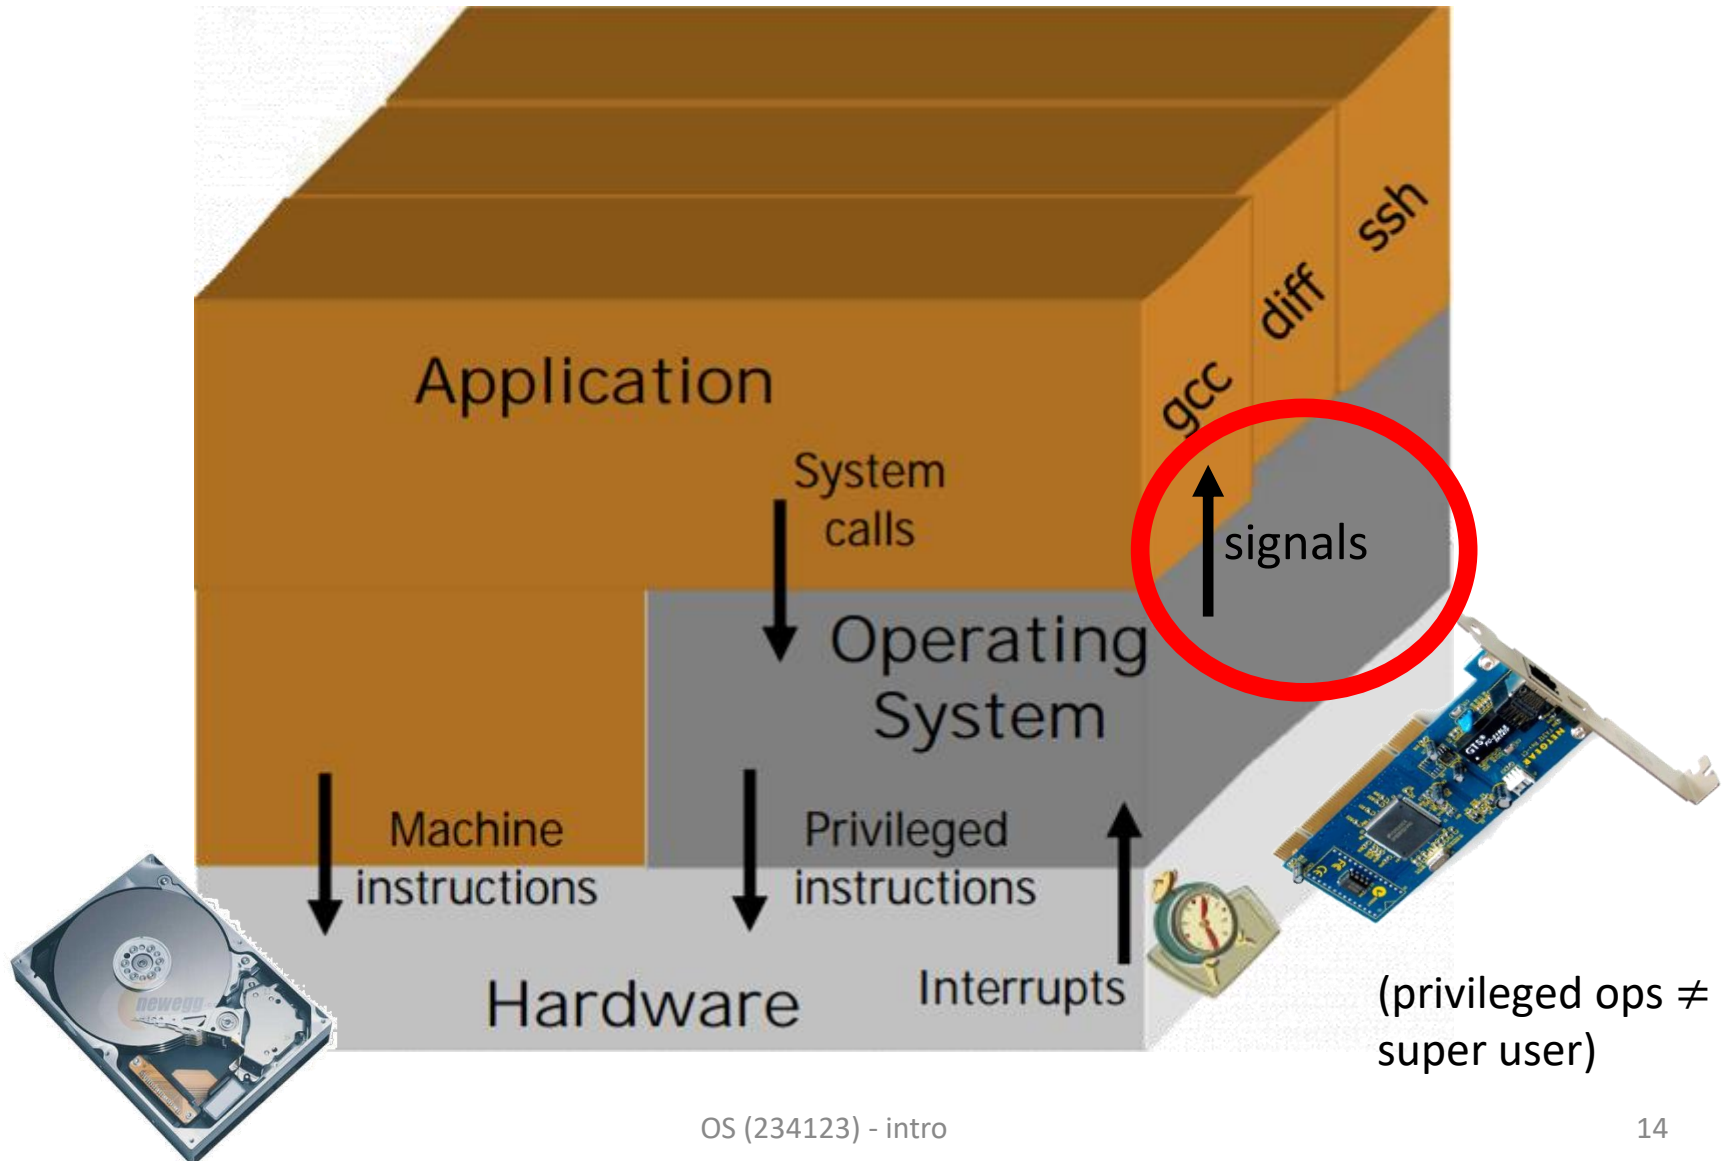
- *“If a parent process terminates, then its zombie children (if any) are adopted by init”*

- <https://linux.die.net/man/2/wait>

OS-supported asynchronous notifications

POSIX SIGNALS

Reminder from the 1st lecture



What are signals & signal handlers

- **Signal = notification “sent” to a process**
 - To asynchronously notify it that some event occurred
- **When receiving a signal**
 - The process stops whatever it is doing & handles it
- **Default signal handling action**
 - Either die or ignore (depends on the type of the signal)
- **The process can configure how it handles most signals**
 - Different signals can have a different handlers,
and they can be temporarily blocked/unblocked = “**masked**”/“**unmasked**”
 - Except for 2 signals (well, actually 3 – discussed shortly)
- **Signals have names and numbers standardized by POSIX**
 - Do ‘man 7 signal’ in shell/Google for a listing/explanation of all signals
 - For example: <http://man7.org/linux/man-pages/man7/signal.7.html> ,
<https://pubs.opengroup.org/onlinepubs/9699919799/basedefs/signal.h.html>
 - **HOMEWORK:** take a few minutes to quickly survey all signals

Silly example

```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>

void sigfpe_handler(int signum) {
    fprintf(stderr, "I divided by zero (sig=%d)!\n",
            signum);    // prints SIGFPE's const value
    exit(EXIT_FAILURE); // what happens if not exiting?
}

int main() {
    signal(SIGFPE, sigfpe_handler);
    int x = 1/0; // processor interrupt, then OS signal
    return 0;
}
```


Another silly example

```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>

void sigint_handler(int signum) {
    printf("I'm disregarding your ctrl-c!\n");
}

int main() {
    // when pressing ctrl-c in the shell
    // => SIGINT is delivered to foreground process
    // (who makes this happen?)
    signal(SIGINT, sigint_handler);
    for(;;) { /*endless loop*/ }
    return 0;
}
```

Another silly example

```
<0>dan@csa:~$ ./a.out
```

```
# here I clicked ctrl-c => delivered SIGINT
```

```
I'm disregarding your ctrl-c!
```

```
# here I clicked ctrl-c again => delivered SIGINT
```

```
I'm disregarding your ctrl-c!
```

```
# here I clicked ctrl-z => delivered SIGSTOP, must obey
```

```
[1]+  Stopped ./a.out
```

```
<148>dan@csa:~$ ps
```

PID	TTY	TIME	CMD
10148	pts/19	00:00:00	bash
21709	pts/19	00:00:12	a.out
21710	pts/19	00:00:00	ps

```
<0>dan@csa:~$ kill -9 21709 # 9=SIGKILL, must obey
```

```
[1]+  Killed ./a.out
```

```
<0>dan@csa:~$
```

Another silly example

```
<0>dan@csa:~$ ./a.out
```

```
# here I clicked ctrl-c (=> deliver SIGINT)
```

```
I'm ignoring your ctrl-c!
```

```
# here I
```

```
I'm ignor
```

```
# here I
```

```
[1]+ Sto
```

```
<148>dan@
```

When I click ctrl-c, the OS gets an interrupt from the keyboard, which the OS then translates into a SIGINT signal delivered to the relevant process.

```
(NT)
```

```
can't ignore
```

```
PID TTY TIME CMD
10148 pts/19 00:00:00 bash
21709 pts/19 00:00:12 a.out
21710 pts/19 00:00:00 ps
```

```
<0>dan@csa:~$ kill -9 21709 # 9=SIGKILL, can't ignore
```

```
[1]+ Killed ./a.out
```

```
<0>dan@csa:~$
```

Notice

- **Argument for the 'kill' shell utility**
 - Any signal (not just 9=kill)
 - kill -9 <pid>
 - kill -s KILL <pid>
 - kill -s SIGKILL <pid>
- **There are 2+1=3 signals that a process can't ignore**
 - SIGKILL = terminate the receiving process
 - SIGSTOP = suspend the receiving process (make it sleep)
- **Is the affect of SIGSTOP reversible?**
 - Yes, when you send to the process SIGCONT
- **SIGCONT can't be ignored either...**
 - A SIGSTOP-ed process **will** continue
 - But process can set a handler for it, which will be invoked immediately when the process gets hit by the SIGCONT and is resumed as a result
- **What can you do with SIGSTOP/SIGCONT?**

Job control

- **In the shell, assume we run a program ‘loop’ that does this**
 - `int main() { while(1); return 0; }`
- **As noted, clicking ctrl-z in the shell**
 - Will make ‘loop’ sleep
- **Subsequently, invoking ‘fg’ in the shell**
 - Will wake ‘loop’ up in the **foreground** (meaning, the shell sleeps, and any typed input is directed to ‘loop’)
- **Alternatively, invoking ‘bg’ in the shell**
 - Will wake ‘loop’ up in the **background** (as is we executed it with “&”, so shell becomes operational, and typed input goes to the shell)
- **Simplifying lie I told**
 - For reasons related to job control, actually, ctrl-z
 - Generates STGTSTP, not SIGSTOP
 - There are subtle differences between the two, which we won’t learn

This is how a signal is truly ignored

```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    // Before, we used sigint_handler as the arg.
    // Now, we use the SIG_IGN macro, which means
    // no handler will be called.
    // There's also SIG_DFL, to restore the default
    // behavior
    signal(SIGINT, SIG_IGN);
    for(;;) { /*endless loop*/ }
    return 0;
}
```

Example – ask a running daemon how much “work” it did thus far



- A “daemon” is
 - Background process, not controlled interactively by user
 - In shell: **nohup** <command> & (see: <https://linux.die.net/man/1/nohup>)
 - Daemon name typically ends with “d” (e.g., sshd, syslogd, swapd)

```
void do_work() { for(int i=0; i<10000000; i++); }
int g_count=0; // counts num. of times do_work was invoked
void sigusr_handler(int signum) {
    printf("Work done so far: %d\n", g_count);
}
int main() {
    signal(SIGUSR1, sigusr_handler);
    for(;;) { do_work(); g_count++; }
    return 0;
}
```

Example – ask a running daemon how much “work” it did thus far



```
<0>dan@csa:~$ ./a.out &
```

```
[1] 23998
```

recall: kill utility also accepts strings as signals

```
<0>dan@csa:~$ kill -s USR1 23998
```

```
Work done so far: 626
```

```
<0>dan@csa:~$ kill -s USR1 23998
```

```
Work done so far: 862
```

```
<0>dan@csa:~$ kill -s USR1 23998
```

```
Work done so far: 1050
```

```
<0>dan@csa:~$ kill -s KILL 23998
```

```
[1]+  Killed                  ./a.out
```


Enumerate signals

- **SIGSEGV, SIGBUS, SIGILL, SIGFPE**
 - **ILL** = illegal instruction (trying to invoke privileged instruction)
 - **SEGV** = segmentation violation (illegal memory ref, e.g., outside an array)
 - **BUS** = dereference invalid address (null/misaligned, assume it's like SEGV)
 - **FPE** = floating point exception (despite name, actually *all* arithmetic errors, not just floating point; example: divide by zero)
 - These are driven by the associated **(HW) interrupts**
 - The OS gets the associated interrupt
 - The OS interrupt handler sees to it that the misbehaving process gets the associated signal
 - The default signal handler for these signals: core dump + die
- **SIGCHLD**
 - Parent gets it whenever fork()-ed child terminates or is SIGSTOP-ed
- **SIGALRM**
 - Get a signal after some specified time
 - Set by system calls: **alarm(2)** & **setitimer(2)** (homework: read man)

Enumerate signals

- **SIGTRAP**
 - When debugging / single-stepping a process
 - E.g., can be delivered upon each instruction
- **SIGUSR1, SIGUSR2**
 - User decides the meaning (e.g., see our daemon example)
- **SIGXCPU**
 - Delivered when a process used up more CPU than its soft-limit allows
 - Soft/hard limits are set by the system call: [setrlimit\(\)](#)
 - Soft-limits warn the process its about to exceed the hard-limit
 - Exceeding the hard-limit => SIGKILL will be delivered
- **SIGPIPE**
 - Write to pipe with no readers (we'll learn about pipes later, for the time being think about the shell's pipe: "|")

Enumerate signals

- **SIGIO**

- Can configure file descriptors such that a signal will be delivered whenever some I/O is ready
- Typically makes sense when also configuring the file descriptors to be “non blocking”
 - E.g., when read()ing from a non-blocking file descriptor, the system call immediately returns to user if there’s currently nothing to read
 - In this case, errno will be set to EAGAIN = EWOULDBLOCK

- **And a few more**

- man 7 signal
- <http://man7.org/linux/man-pages/man7/signal.7.html>

Signal vs. interrupts

	interrupts	signals
Who triggers them? Who defines their meaning?	Hardware: CPU cores (sync) & other devices (async)	Software (OS), HW is unaware
Who handles them? Who (un)blocks them?	OS	processes
When do they occur?	Both synchronously & asynchronously	Likewise, but, technically, invoked when returning from kernel to user

Signal system calls – sending

- **int kill(pid_t *pid*, int *sig*)**
 - (Not the shell utility, the actual system call)
 - Allows a process to send a signal to another process (or to itself)
 - Homework: How?
 - man 2 kill – <http://linux.die.net/man/2/kill>
 - “2” is for system calls
 - “1” is for shell utilities

Signal system calls – (un)blocking

- **Signals can be asynchronous => might lead to "race conditions"**
 - Therefore, as noted, all signals (except kill/stop) can be blocked
 - Like how OS disables/enables interrupt
- **How**
 - The PCB maintains a set of currently blocked signals
 - Which can be manipulated by users via the following syscall
- **int sigprocmask(int *how*, const sigset_t **set*, sigset_t **oldset*)**
 - 'how' = SIG_BLOCK (+=), SIG_UNBLOCK (-=), SIG_SETMASK (=)
 - 'set' can be manipulated with sigset ops [sigemptyset](#)(sigset_t *set), [sigfillset](#)(sigset_t *), [sigaddset](#)(sigset_t *set, int signum), [sigismember](#)(sigset_t *set, int signum)
- **Manual**
 - man 2 sigprocmask – <http://linux.die.net/man/2/sigprocmask>
 - man 3 sigsetops – <https://linux.die.net/man/3/sigsetops>

(Un)blocking example

```
Record_t db[N];

void my_handler(int signum) {
    /* may read/update db */
}

int main(int argc, char *argv[])
{
    sigset_t mask, orig_mask;
    sigemptyset(&mask);
    sigaddset(&mask, SIGTERM);

    signal(SIGTERM, my_handler);

    for(;;) {
        char *cmd = read_command();
        sigprocmask(SIG_BLOCK, &mask, &orig_mask);
        // do stuff that may read/update db...
        sigprocmask(SIG_SETMASK, &orig_mask, NULL);
    }
    return 0;
}
```

Recall that every syscall might fail;
the example ignores this for brevity

Signal system calls – control & more info

- **Additional info about signal & fine-grain control**
 - over how signals operate is provided via the following syscall
- **int sigaction(int signum,
 const struct sigaction *act,
 struct sigaction *oldact)**
 - man 2 sigaction – <http://linux.die.net/man/2/sigaction>
(homework: read it)

Signals interact with other system calls

- **ssize_t read(int fd, void *buf, size_t count);**
 - What happens if getting signal while read()ing?
 - The read system call returns -1, and it sets the global variable 'errno' to hold EINTR
 - An example whereby read() might fail and user should simply retry

```
int readn( int sockfd /*learn later*/, char *ptr, int nbytes )
{
    int nleft = nbytes;

    while( nleft > 0 ) { // 'read' is typically done in a loop (why?)
        int nread = read(sockfd, ptr, nleft);
        if( (nread == -1) && (errno != EINTR) ) {
            fprintf(stderr, "read failed, errno=%m", errno);
            return -1;
        }
        else if( nread == 0 )
            break; /*EOF*/

        nleft -= nread;
        ptr   += nread;
    }
    return nbytes - nleft;
}
```